



Octopus User Manual

Version 0.3 alpha
31 October 2017
University of Oxford

OVERVIEW	5
Introduction	5
What's in this manual?	5
License and copyright	5
Technical assistance	5
INTRODUCTION	6
Variant calling	6
Hybrid mapping based variant calling	6
Haplotype based variant calling.....	6
Local call phasing	6
CALLING MODELS	7
Individual	7
Population	7
Trio	7
Cancer	7
INSTALLATION	8
System Requirements	8
Hardware	8
Required Software	8
Optional software	9
Downloading.....	9
Building	9
Easy install with Python	9
Building with CMake	9
Debug builds	10
RUNNING TESTS	10
GETTING STARTED	10
Basic usage.....	10
Required arguments	10
Optional arguments	11
Reporting bugs.....	11
Requesting new features.....	11
BEST PRACTICES	12

Reference selection	12
Read mapping	12
Read preprocessing	12
Variant calling	12
Variant call filtering	12
COMMAND LINE REFERENCE	13
General	13
Read Pre-Processing	15
Variant Generation	17
Haplotype Generation	18
Calling	19
Trio	19
Cancer	20
Phasing	20
Advanced	20
Call Filtering	21
PERFORMANCE OPTIMISATION	22
Execution time	22
Memory consumption	22
Multithreading	23
Variant generation	23
Haplotype generation and phasing	23
Calling model selection and parametrisation	24
Variant filtering	24
EXAMPLES	26
Calling germline variants in a single sample	26
Calling variants in a targeted exome panel	26
Ignoring decoy contigs from a whole genome run	26
Calling germline variants in a population (in development)	27
Calling de novo mutations in a trio	27
Calling somatic mutations in a tumour-normal pair	27
HLA genotyping	27
Calling variants in haploid organism	27
Running in multithread mode	28

Using a configuration file	28
TROUBLESHOOTING	29
Building	29
Why are the requirements so strict?	29
CMake chooses a bad compiler	29
Compilation fails	29
Linking fails	29
Boost libraries fail to link	29
Compilation has lots of #pragma warnings	30
Runtime	30
Segmentation fault.....	30
Execution is slow	30
Execution delays after initialising calling components in threaded mode.....	30
Run hangs in decoy contigs	30
Behaviour	31
No calls are reported	31
Regions are skipped because of too many haplotypes	31
A call changes when a different input region is given	31
Why doesn't octopus report genotype likelihoods?	31
Why do octopus VCF files contain * and .?	31
SNP accuracy improves in fast mode	32
Calling performance is worse with assembler.....	32
CONTACT	33
APPENDIX	33
Installing Requirements.....	33
OS X	33
Ubuntu	33
Variant Generation	34
Haplotype Generation.....	34
Phasing	34
Variant Call Filtering.....	35
GLOSSARY	35

OVERVIEW

Introduction

Octopus is a command line tool that detects **genetic variation** from **high-throughput sequencing data** (reads) relative to a **reference sequence**. The tool must be provided with an indexed FASTA reference file and one or more SAM format mapped and aligned read files, and will produce a VCF format file containing phased variant calls. Octopus can be used detect and classify germline, somatic, or *de novo* mutations across multiple samples.

What's in this manual?

This is a **user** manual intended for novice to advanced users to get octopus running optimally. It is not intended to give a detailed description of the algorithms implemented in octopus (although some pertinent details are given in the appendix to help users understand some parameters), nor is it a technical manual for software developers. Please refer to the [octopus paper](#) and [developer manual](#) for detailed descriptions of these topics.

License and copyright

Octopus is distributed under the [MIT license](#). The copyright holder is the Daniel Cooke (daniel.cooke@well.ox.ac.uk) who reserve the right to change the license terms.

Technical assistance

- For assistance with specific technical or conceptual issues not addressed in this manual, please contact daniel.cooke@well.ox.ac.uk.
- For general discussion, please use the [octopus Gitter chat](#).
- For bugs and feature requests, please use the [octopus issue tracker](#).

INTRODUCTION

Variant calling

Variant calling is an inference task; the aim is to report the underlying genome of the sample under consideration, given a set of indirect observations of the sample's genome (reads). This is a statistical problem as the underlying data is noisy due to the sequencing process (errors can be introduced in the library preparation and sequencing itself). In practice, not all inferred genetic information is reported as the vast majority of information is conserved amongst populations. Instead only *differences* compared to a reference sequence for the population are reported, these are called variants.

Hybrid mapping based variant calling

Mapping based variant callers require preprocessed input from a *read mapper*. A read mapper takes raw sequencing reads and attempts to determine the origin of each read *independently* relative to a reference genome. Most mappers will subsequently align the read around the mapped location.

The variant calling task is much simplified when read mapping information is available as the domain of possible variants is significantly reduced. However, with mapping based approaches, the overall accuracy of the caller may be bounded by the accuracy of the mapper and alignment algorithm; the mapping stage itself can be viewed as a variant calling process as the mapper must also account for deviations from the reference sequence due to real variation and sequencing errors.

The other method is to avoid using a mapper entirely; just take the raw reads and assemble them into full contigs. Such approaches do exist, and usually employ De Bruijn graphs and similar algorithms, but are usually underpowered compared to mapping based approaches. There is also a significant computational overhead attached to assembly based approaches.

Experience is showing that the best overall solution is a hybrid approach where read mapping information is used, but only partially. Reads mapping within a certain genomic interval are locally reassembled and *then* aligned to the assembled contig. The idea being that the read mapper may be wrong, but it is unlikely to be very wrong; the true read origin is unlikely to be far away from the mapped location.

Haplotype based variant calling

Haplotype based variant callers attempt to jointly genotype more than one genomic position simultaneously. This is in contrast to traditional *positional* based variant calling that only genotypes a single location at once. The advantage of haplotype based is that the space of possible *errors* increases exponentially with haplotype length, while the space of true haplotypes remains constant in the number of samples and organism ploidy. It is therefore much easier to classify true variation and sequencing errors.

Local call phasing

Phasing refers to assigning called alleles to a particular haplotype; calls are *phased* if information indicating which called alleles occur on the same haplotype is provided. It is only possible to fully reconstruct a sample's genome if phased calls are generated. Octopus is able to generate phased calls - the phase information is provided in the final call set.

CALLING MODELS

Octopus is really a framework for genotyping samples given different states of knowledge about those samples, such as different sample biology or experimental method, expressed via a **calling model**. A calling model serves two purposes: firstly, it defines the type of calls and inferences that should be made (e.g. *de novo* classification), and secondly, it defines a probability model to calculate posterior probabilities for the given call types. Octopus currently has four calling models, which are briefly discussed below.

Individual

The individual calling model is the simplest, it is intended to model a single healthy individual with known chromosome copy number (ploidy) for all chromosomes (or contigs). The advantage of having a bespoke model for an individual is that the genotype posterior distribution can be calculated exactly.

Population

The population calling model is intended for genotyping multiple unrelated samples from a population. Like the individual model, it is assumed each sample is healthy with a known chromosome copy number. The first advantage of calling samples jointly, as apposed to calling each sample individually and then merging the results, is that power is increased to call common variation. This is particularly true for low coverage data. The second advantage is that genotyping samples jointly allows a consistent call set to be produced; merging independent call-sets can be very challenging.

Trio

The trio calling model is used to genotype a family consisting of a mother, father, and offspring. All members of the trio are assumed to be healthy with known chromosome copy numbers. However, unlike the population model, the trio model explicitly models the relationship between the samples, and is therefore able to classify *de novo* mutations in the child.

Cancer

The cancer model is used to genotype tumours from a single individual. All tumours are assumed to be metastasis from the same primary tumour (or the primary itself). The model can be used to classify somatic mutations, and infer local copy number changes around called mutations. Unlike the other calling models, the chromosome copy number of each tumour is not assumed to be known, however, if a normal sample with known chromosome copy number is also present the classification power of the model is increased.

INSTALLATION

This section gives detailed instructions on how to obtain, build, and install octopus. Please refer to the troubleshooting section for common installation problems not addressed here.

SYSTEM REQUIREMENTS

Octopus is mostly written in C++, and therefore requires the source code to be compiled for the target machine architecture with a C++ compiler. You will need to consult your operating systems technical documentation to determine a suitable C++ compiler.

Hardware

In principle octopus can run on any machine capable of compiling a C++14 program. However, given the complex numerical algorithms involved in running octopus the following guidelines are offered¹:

Technology	Minimum	Recommended
Processor	Intel Core i5	32 x Intel Core i5
Memory	8GB	16GB
Disk	500GB	1TB

Most modern desktop and laptop computers should satisfy these requirements. The user should understand that hardware requirements will vary greatly depending on the use-case and workload. For example, calling many high coverage samples will require far greater memory than a single low coverage sample. Octopus is fully multithreaded, and to achieve reasonable runtime performance on large tasks it is recommended to make multiple processor cores available.

Octopus requires SSE2 hardware support.

Required Software

- A C++14 compliant compiler with SSE2 support²
- An implementation of the C++14 standard library³
- Boost 1.61 or greater
- htlib 1.3.1 or greater
- CMake 3.5 or greater

¹ These guidelines are based on running octopus on a single high coverage (~50x) human sample.

² GCC 6.1.1 and below have bugs which affect octopus; only use GCC 6.2 and above. LLVM Clang 3.8 has been tested and compiles. Visual Studios and Intel C++ compilers have not been tested.

³ It is highly recommended to use the compilers native C++ standard library implementation: libstdc++ for GCC and libc++ for Clang.

Optional software

- Git 2.5 or greater
- Python3

Instructions on obtaining the requirements on OS X and Ubuntu are given in the Appendix.

DOWNLOADING

Octopus is distributed via the project hosting website Github. There are two ways to obtain a copy of the source code from Github:

1. Visit the [octopus Github webpage](#) and click the *Clone or download* box. This will download a zip file named **octopus-master.zip** containing the octopus source code. Move the zip file to a suitable location, unzip it, and rename the folder to **octopus**.
2. Open a command line terminal and move to a directory where you would like octopus to be downloaded, then execute the git command:

```
$ git clone https://github.com/luntergroup/octopus.git
```

The octopus source code will be downloaded into a folder named **octopus**.

BUILDING

Once the source code is obtained there are two methods to create an executable for your target machine, both require *CMake* to generate a native makefile which is used by a native build-tool to build the final executable. It is highly recommended to do an out of source build.

Easy install with Python

In the top level directory there is a Python3 script `install.py` which will execute all the necessary build steps:

```
$ ./install.py
```

Which will install into the octopus `bin` directory. To install into a standard location use:

```
$ ./install.py --root
```

If *CMake* is not able to find a suitable C++ compiler, it may be necessary to explicitly specify where such a compiler exists:

```
$ ./install.py --compiler /path/to/compiler/cpp
```

Building with CMake

It is also possible to build the source directly with *CMake*:

```
$ cd build
```

```
$ cmake ..  
$ make install  
$ cmake -D INSTALL_ROOT=ON ..  
$ cmake -D CMAKE_CXX_COMPILER=g++-4.2 ..
```

Using the python script is recommended however as it ensures an out of source build.

Debug builds

It is possible to build octopus with debug information. This is only recommended for debugging and will hopefully not be needed for users. To do so, simply add the command `--debug` to the Python install script.

RUNNING TESTS

If you downloaded a developmental version of octopus, it is good practise to run all the packaged tests before using any of the tools for production work. The release versions are guaranteed to have passed all tests. To install octopus for testing and run the tests use:

```
$ test/install.py
```

Like the other install script this command can also be supplied with a compiler.

GETTING STARTED

Once successfully installed octopus is ready to run. This section is for novice users who want a gentle introduction to variant calling. Advanced users should consult the command line reference section for detailed descriptions of specific features.

Basic usage

Octopus is a command line tool and must be executed from a command terminal. The simplest octopus run is without any arguments:

```
$ octopus
```

Which will report a user error informing there are missing required arguments! You can request a reminder of all required and optional parameters with the `--help` command:

```
$ octopus --help
```

This will display a similar table to the command reference below.

Required arguments

Only two command line argument are required. First, the reference genome to use for analysis specified with `--reference; -R`, which must be given a path to a FASTA file containing the reference genome. A FASTA index file with the same name, but extension `.fai` is also required to exist in the same directory as the given reference.

Second, a list of read file (BAM or CRAM format) paths must be supplied. These can either be supplied directly with the `--reads; -I` option, or with the `--reads-file; -i` option, which must be given a path which itself contains a list of paths to read files. These two options can also be used conjunctively, any duplicate files will be ignored. Each read file must have an associated index file that exists in the same directory as the read file (`.bai` for BAM and `.crai` for CRAM).

Optional arguments

Octopus has many optional arguments that affect accuracy and runtime performance. The default parameters have been chosen with human germline sequence data mapped with BWA-MEM in mind; many users will find the default arguments offer adequate performance on human samples. For non human data samples, the default parameters may not offer good performance, especially for non-diploid organisms, and users are advised to carefully read the available options. Even for users only interested in human samples, it is recommended they briefly acquaint themselves with the available options. A detailed description of all command line options can be found later in this manual.

Reporting bugs

While much effort has gone into making octopus stable and bug free, it is possible that some remain. Please check the troubleshooting pages before reporting potential bugs, and it may also be beneficial to first report the issue on the forums - someone else may have had the same issue!

Once a bug has been verified, try if possible to find a minimal verifiable example (MVE); that is the least amount of data (usually this involves selecting smaller target regions) that triggers the problematic behaviour. Once an MVE has been found, if the problem is a reported exception, use the `--debug` command to generate a log file that can be sent to maintainers. If the problem is a segmentation fault or similar, recompile octopus in debug mode (see install instructions) and rerun to see if any assertions are violated. If no assertion is violated and the user is able, it would be very helpful to run `valgrind` on the MVE and send the report to maintenance. Please note that the log file cannot be used to resolve all issues; maintainers may require access to raw data to fix some issues, it is therefore vital to have a MVE. Bugs can be reported to the [octopus issue tracker](#).

Requesting new features

Feedback is very welcome! Please start by suggesting a new feature on the [octopus forum](#) and then if well received make an official feature requests to the [octopus issue tracker](#).

BEST PRACTICES

This section gives some brief advice on best practise workflow from FASTQ to VCF.

Reference selection

Use the latest possible reference genome for your sample. Reference assemblies are often updated to reflect resolutions of complex loci, or to add decoy sequence which reduces mapping issues and improve calling quality.

Read mapping

Octopus requires mapped and aligned reads in the SAM format. The quality of the mapping software is therefore an essential part of the variant calling process. While the performance of mappers can vary considerably depending on the type of sequencing data used, BWA-MEM (default settings) is recommended as it is widely used, well tested, and has been shown to perform well on a wide range of data - in particular human genetic data.

Read preprocessing

Octopus does not require any read preprocessing after mapping, such as duplicate marking, indel realignment, or base quality recalibration. Unlike other variant callers, octopus is unlikely to benefit from such techniques as reads are preprocessed internally, indels are essentially realigned during calling, and base quality scores are also internally manipulated depending on sequence context. However, if your data is already preprocessed, octopus should perform equally well with this data.

Variant calling

Ensure the correct calling model is selected for the type of data to be analysed. Look at the private parameters for the chosen calling model and verify the defaults are reasonable. At the very least, check the ploidy assumptions are correct. Once the calling model is appropriately configured, consult the performance optimisation section to help tune other calling parameters.

Variant call filtering

This version of octopus provides threshold based variant call filtering. We recommend using the default filter expression which is intended to exclude high quality false positives. Variant call filtering is a difficult problem (see appendix), in particular there are numerous problems with threshold based filtering. We are actively working on improved filtering methods for new releases.

COMMAND LINE REFERENCE

This section contains a description of each command line option. The commands are separated into sections which roughly correspond to different area of concern. At the end of each section a detailed explanation of any non-trivial commands is given.

Some options have so called default *implicit* values, that is, they are by default disabled, but can be enabled with the implicit default value by just specifying the option name. To override the default implicit value, you must use the form `--option-name=value`; the form `--option-name value` will **not** work! Implicit options are labelled as `=(default implicit value)`.

Entries with a red border are currently placeholders and are not yet implemented, they are included to give an indication of what will be available in the first official release. Entries with an orange border are currently implemented, but are likely to change before the first official release.

GENERAL

Command	Description	Default value
<code>--reference, -R</code>	The reference genome to use for analysis. Must match the reference genome used to map reads against.	None
<code>--reads, -I</code>	The read files to use for analysis. Can be specified multiple times and given a space separated list of argument.	None
<code>--reads-file, -i</code>	A path that contains a list of read file paths, one per line, to use for analysis.	None
<code>--regions, -T</code>	A list of genomic intervals to analyse.	All regions present in the reference index.
<code>--regions-file, -t</code>	A path to a file that contains a list of genomic intervals to analyse. Must have one region per line. BED format is accepted.	None
<code>--skip-regions, -K</code>	A list of genomic intervals that should be ignored.	None

Command	Description	Default value
<code>--skip-regions-file, k</code>	A path to a file that contains a list of genomic intervals that should be ignored. Must have one region per line. BED format is accepted.	None
<code>--one-based-indexing</code>	Reads all user input regions using one-based indexing rather than zero based.	Disabled
<code>--samples, -S</code>	A list of samples to analyse, which must be a subset of those samples in the reads.	All samples found in the reads.
<code>--samples-file, -s</code>	A path to a file containing a list of samples to analyse.	None
<code>--pedigree</code>	PED file containing sample pedigree. Only currently used by trio calling model.	None
<code>--fast</code>	Disables various algorithmic features to significantly reduce runtime, at the cost of worse calling accuracy. Equivalent to <code>-a off -l minimal -x 50</code> .	Disabled
<code>--very-fast</code>	Disables various algorithmic features to significantly reduce runtime, at the cost of worse calling accuracy. Equivalent to <code>--fast --inactive-flank-scoring off</code> .	Disabled
<code>--threads</code>	Enables multithreading. If not supplied with an argument (recommended), the number of threads is automatically determined. Otherwise the number of threads is limited to the given number.	Disabled. =(automatic)
<code>--working-directory, -w</code>	Any path given to octopus will be relative to the working directory, unless the path is already valid.	None
<code>--max-reference-cache-footprint, -X</code>	The maximum amount of memory available to cache reference sequence. Caching reference sequence reduces file IO.	500MB
<code>--target-read-buffer-footprint, -B</code>	The recommended amount of memory available for buffering read data. This is not a strict limit.	2GB
<code>--max-open-read-files</code>	Limits the number of open read files to the given number. Note each read file also has an index which is not accounted for.	250
<code>--contig-output-order</code>	Which order should contigs appear in the final output? Possible values are: <code>lexicographicalAscending</code> , <code>lexicographicalDescending</code> , <code>contigSizeAscending</code> , <code>contigSizeDescending</code> , <code>asInReference</code> , <code>asInReferenceReversed</code> .	<code>asInReference</code> ndex
<code>--sites-only</code>	Remove genotype calls and associated information from final VCF output.	Disabled

Command	Description	Default value
<code>--regenotype</code>	A VCF file specifying sites to regenotype; only calls listed in this files will appear in the output.	None
<code>--legacy</code>	Outputs a more conventional VCF file in addition to the standard octopus format.	Disabled
<code>--debug</code>	Writes verbose debug information to a log file. Can be supplied with a path.	Disabled. =(octopus_debug.log)
<code>--trace</code>	Writes very verbose debug information to a log file. For maintainer use only. Can be supplied with a path.	Disabled. =(octopus_trace.log)
<code>--version</code>	Displays the current version number and other meta information.	None
<code>--config</code>	A configuration file that contains values for some or all of the options listed here.	None

READ PRE-PROCESSING

Command	Description	Default value
<code>--read-transforms</code>	Use to turn off all read transformations. Reads can still be filtered.	Enabled
<code>--soft-clip-masking</code>	Use to turn off soft clip masking (assigning base quality zero) of soft clipped read flanks.	Enabled
<code>--mask-low-quality-tails</code>	Masks (assigns base quality zero) the tail given number of bases of each read.	Disabled. =(3)
<code>--mask-soft-clipped-boundries</code>	Masks (assigns base quality zero) to the soft clipped flanks of reads, plus an additional number of given bases.	Disabled
<code>--adapter-masking</code>	Prevents read bases that are considered likely adapter contaminants, as determined by octopuses native adapter contamination detector, from being masked (assigned base quality zero). This command is redundant unless the command <code>--allow-adapter-contaminated-reads</code> is also used.	Enabled
<code>--overlap-masking</code>	Prevents masking (assigning base quality zero) of read bases that overlap (w.r.t mapping location) of other segments within the reads template. For paired-end reads, this usually refers to the reads mate. Only one corresponding base of each read is masked; the other is left untouched.	Enabled

Command	Description	Default value
<code>--read-filtering</code>	Prevents any read from being quality control filtered, this does not affect downsampling.	Enabled
<code>--consider-unmapped-reads</code>	Turns off filtering of reads marked as unmapped. Note this is not the same as reads with mapping quality zero.	Disabled
<code>--min-mapping-quality</code>	Discards reads with mapping quality less than this before calling.	20
<code>--good-base-quality</code>	The base quality threshold to use for the the options <code>--min-good-base-fraction</code> and <code>--min-good-bases</code> .	20
<code>--min-good-base-fraction</code>	The maximum fraction of bases below <code>--min-good-base-quality</code> before the read is discarded.	Disabled
<code>--min-good-bases</code>	The minimum number of bases equal to or above <code>--min-good-base-quality</code> before a read is considered.	20
<code>--allow-qc-fails</code>	Prevents removal of reads marked as QC failed.	Disabled
<code>--min-read-length</code>	Discards reads with less bases than this.	None
<code>--max-read-length</code>	Discards reads with more bases than this.	None
<code>--allow-marked-duplicates</code>	Prevents removal of reads pre-marked as duplicates.	Disabled
<code>--allow-octopus-duplicates</code>	Prevents removal of reads that octopuses native duplicate detector marks as duplicates.	Disabled
<code>--no-secondary-alignments</code>	Discard reads marked as being secondary alignments.	Allow
<code>--no-supplementary-alignments</code>	Discards reads marked as being supplementary alignments.	Allow
<code>--consider-reads-with-unmapped-segments</code>	Prevents removal of reads where one or more segments in the reads template are marked as unmapped. For paired-end reads, this usually refers to the read mate.	Disabled
<code>--consider-reads-with-distance-segments</code>	Prevents removal of reads that have template segments mapped to a different contig. For paired end reads, this usually refers to the read mate.	Disabled
<code>--no-adapter-contaminated-reads</code>	Prevents removal of reads that are likely to contain adapter contamination, as determined by octopuses native adapter contamination detector.	Disabled
<code>--disable-downsampling</code>	Turns off all downsampling. Reads may still be filtered.	Enabled

Command	Description	Default value
<code>--downsample-above</code>	Trigger downsampling of a sample when the read depth in a region is above this value.	500
<code>--downsample-target</code>	Once a region has been flagged for downsampling, try to remove reads in the region to achieve this level of coverage. Must be greater than <code>--downsample-above</code> .	400

VARIANT GENERATION

Command	Description	Default value
<code>--raw-cigar-candidate-generator, -g</code>	Do not use the raw cigar variant candidate generator to propose candidate variants.	Enabled
<code>--assembly-candidate-generator, -a</code>	Do not use the local reassembler generator to propose candidate variants.	Enabled
<code>--source-candidates</code>	Consider all sites in the given VCF format as candidate variants. This differs from the option <code>--regenotype</code> as the final call set is not required to be a subset of these calls.	None
<code>--max-variant-size</code>	The maximum variant size (w.r.t genomic interval span) that any candidate variant generator may propose.	2,000
<code>--min-supporting-reads</code>	Overrides the default raw cigar generator and applies a simple threshold inclusion predicate based on the number of observed reads. Observations must have base quality greater than that indicated in <code>--min-base-quality</code> .	2
<code>--min-base-quality</code>	The minimum base quality a read base must have before it is considered as supporting a variant.	20
<code>--kmer-sizes</code>	Default k-mer sizes to use by the assembler. If none of these sizes are successful then 6 more fallback sizes, in increments of 10 past the largest size given, will be tried.	10 20
<code>--num-fallback-kmers</code>	The number of fallback k-mer sizes to try if the default ones fail to provide a valid graph.	7
<code>--fallback-kmer-gap</code>	The gap size of fallback k-mers.	10
<code>--max-region-to-assemble</code>	The maximum region size that will be used for local reassembly. Larger sizes may result in larger structural variation being found, but reduces sensitivity to smaller variation.	400

Command	Description	Default value
<code>--max-assemble-region-overlap</code>	The maximum number of bases assembly windows are allowed to overlap. A higher overlap may increase sensitivity but increase runtime.	200
<code>--force-assemble</code>	Forces local reassembly of all genomic regions.	FALSE
<code>--assembler-mask-base-quality</code>	Mismatching bases with quality less than this will be masked as reference before being threaded into the assembly graph.	10
<code>--min-kmer-prune</code>	The minimum number of k-mer observations to keep the k-mer in the graph after pruning.	2
<code>--max-bubbles</code>	The maximum number of bubbles to extract from the assembly graph.	30
<code>--min-bubble-score</code>	The minimum bubble score to extract from the assembly graph.	2

HAPLOTYPE GENERATION

Command	Description	Default value
<code>--max-haplotypes, -x</code>	The maximum number of haplotypes that can be used to generate candidate genotypes. If the haplotype generator proposes more haplotypes than this then the excess will be filtered.	200
<code>--haplotype-holdout-threshold</code>	If a region contains more haplotypes than this, then a subset of alternative alleles will be temporarily removed (held out) and only be analysed once some haplotypes have been discarded.	2,500
<code>--haplotype-overflow</code>	The maximum number of haplotypes a region may have before the region is unconditionally skipped (without attempting to hold out alternative alleles).	200,000
<code>--max-holdout-depth</code>	The maximum number attempts to hold out alternative alleles in a region before the region is skipped.	20
<code>--extension-level</code>	Level of haplotype extension. Possible values are conservative, normal, optimistic, and aggressive.	Normal

The option `--max-haplotype` is a target for the haplotype generator as well as a strict limit for the caller; the haplotype generator will attempt to satisfy the request, but if it fails to do so, the caller will filter the generated haplotype set to this limit.

CALLING

Command	Description	Default value
<code>--caller, -C</code>	Which calling model to use.	individual or population
<code>--organism-ploidy, -P</code>	The autosome ploidy for the analysed organism. All contigs will have this ploidy unless marked otherwise.	2
<code>--contig-ploidies, -p</code>	Assigns ploidies to contigs, overriding the default organism ploidy.	MT=1 Y=1
<code>--snp-heterozygosity</code>	The SNP heterozygosity in the sample population.	0.001
<code>--snp-heterozygosity-stdev</code>	The SNP heterozygosity standard deviation in the sample population.	0.01
<code>--indel-heterozygosity</code>	The INDEL heterozygosity in the sample population.	0.0001
<code>--min-variant-posterior</code>	The minimum posterior probability (QUAL) for a variant to be reported.	2
<code>--use-uniform-genotype-priors</code>	Use uniform genotype priors.	Disabled
<code>--model-posterior</code>	Calculate model posteriors for every call.	Disabled
<code>--refcall</code>	Report reference confidence calls.	Disabled
<code>--min-refcall-posterior</code>	The minimum posterior probability (QUAL) for a reference allele to be reported.	2

TRIO

Command	Description	Default value
<code>--maternal-sample; -M</code>	Which of the given samples is the mother in the trio.	None
<code>--paternal-sample; -F</code>	Which of the given samples is the father in the trio.	None
<code>--snv-denovo-mutation-rate</code>	The germline snv <i>de novo</i> mutation rate.	10 ⁻⁹
<code>--indel-denovo-mutation-rate</code>	The germline indel <i>de novo</i> mutation rate.	10 ⁻¹⁰
<code>--min-denovo-posterior</code>	The minimum posterior probability (phred scale) to emit a <i>de novo</i> mutation call.	0.5
<code>--denovos-only</code>	Only report variant calls classified as <i>de novo</i> mutations. Default is to report both germline and <i>de novo</i> calls.	Disabled

CANCER

Command	Description	Default value
<code>--normal-sample; -N</code>	Which of the given samples is the normal.	None
<code>--somatic-mutation-rate</code>	The somatic mutation rate for the cancer to be analysed.	10 ⁻⁵
<code>--min-expected-somatic-frequency</code>	The minimum expected somatic allele frequency in the sample.	0.05
<code>--min-credible-somatic-frequency</code>	The minimum inferred somatic allele frequency that will be emitted.	0.01
<code>--credible-mass</code>	Mass of the posterior allele frequency distribution to use when calculating allele frequency.	0.99
<code>--min-somatic-posterior</code>	The minimum posterior probability an allele is somatic to be reported.	0.5
<code>--max-cancer-genotypes</code>	The maximum number of cancer genotypes to consider when modelling somatic mutations.	20,000
<code>--somatics-only</code>	Only report variant calls classified as somatic mutations. Default is to report both germline and somatic calls.	Disabled

PHASING

Command	Description	Default value
<code>--phasing-level, -l</code>	The level of phasing. Possible values are: minimal, conservative, moderate, normal, and aggressive.	Normal
<code>--min-phase-score</code>	The minimum phase score (phred scale) a potential phase set may have to be called.	20
<code>--use-unconditional-phase-score</code>	Use the full genotype posterior distribution when calculating phase score, rather than conditioning on called genotypes.	Disabled

ADVANCED

Command	Description	Default value
<code>--haplotype-extension-threshold, -e</code>	Haplotypes with posterior probability (of occurrence in the sample set) can be removed before haplotype extension.	100
<code>--inactive-flank-scoring</code>	Use to disable calculation to account for flank mismatches in HMM routine.	Enabled

Command	Description	Default value
<code>--model-mapping-quality</code>	Use read mapping quality in read likelihood calculation.	Enabled
<code>--max-joint-genotypes</code>	Maximum number of joint genotype vectors that can be considered (applicable to population and trio calling models).	1,000,000
<code>--sequence-error-model</code>	The sequencing error model to use for read likelihood calculation. Possible values are <i>hiseq</i> and <i>x10</i> .	<i>hiseq</i>

CALL FILTERING

Command	Description	Default value
<code>--call-filtering, -f</code>	Use to enable Call Set Refinement (CSR).	Enabled
<code>--filter-expression</code>	Boolean expression to use to filter calls. Current version only supports OR operations and the measure name must appear on the left hand side of the comparator.	QUAL < 10 MQ < 30 MP < 20 AF < 0.05 SB > 0.9 MQD > 0.5
<code>--use-calling-reads-for-filtering</code>	Use the reads used for calling for filtering. Otherwise filtering reads will use default read filters and transforms.	Disabled

PERFORMANCE OPTIMISATION

Octopus is a sophisticated program with many parameters. The default values for these parameters have been chosen with an emphasis on accuracy, while keeping runtime reasonable; they should provide good all round performance for most users. However, if octopus is not performing adequately on your particular dataset it may be due to non-optimal parameterisation. Generally, there is a direct tradeoff between calling accuracy and runtime resource consumption (memory and CPU time).

Execution time

By default octopus favours slower, more accurate variant calling. If accuracy is not critical, in particular around complex regions, it is possible to achieve significant reductions in runtime by altering the behaviour of certain components. In summary, the main components of interest are:

Component	Associated commands	Explanation
Variant generation	<code>-a, --kmer-size</code>	The number of candidate variants that must be considered directly affects runtime complexity. In general, more sensitive variant generation will result in more accurate results but longer runtimes. However, it is important to be aware generating too many candidate variants may actually decrease accuracy as the model may become overwhelmed. Variant generation using local reassembly is also inherently computationally expensive, and the proportion of sites that must be assembled will directly affect runtime.
Haplotype generation	<code>-x</code>	The number of haplotypes considered has a direct impact on calling accuracy and runtime.
Phasing	<code>-l</code>	Phasing requires potentially multiple marginalisation across the entire genotype posterior distribution. This can be costly when there are many genotypes, or many candidate variants.
Calling model	Model dependent (see below)	Some calling models have computationally complex inference procedures. For example, the cancer calling model implements an iterative Variational Bayes algorithm. These algorithms usually have convergence and performance limits which can affect runtime.

There are two convenience command line options `--fast` and `--very-fast` that can be used that automatically adjust these parameters to achieve exceptional runtimes.

Memory consumption

Memory consumption will naturally fluctuate during a run depending on the complexity of the region currently being analysed, however, by far the main source of memory consumption in octopus is from buffering read data. While the size of the read buffer is not directly controllable, the user is able to hint at the maximum buffer

size with the `--target-read-buffer-footprint` option. As this is only a hint, it can be ignored, but in most cases the request is satisfied. The exception to this is when multithreading is used, where there is a greater possibility of the requested limit being exceeded. In this case it is recommended to set the target around 20% less than the true limit.

Another source of memory consumption is reference caching which can be used to improve runtime execution speed as it reduces the need to read from disk. A moderate default reference cache size is used, which can be adjusted by the user.

Finally, significant memory usage can be observed if unreasonable calling parameters are chosen. For example, setting `--max-haplotypes` above 5,000 is likely to lead to memory explosion. Octopus may issue a warning in such cases, but will try to satisfy the user's request.

Multithreading

The default behaviour is to run using only a single thread. To enable octopuses built in multithreading capabilities, the `--threads` command must be used. If an argument is provided to this command then octopus will use this many threads, if no argument is given to this command then octopus will automatically determine the number of threads to use. It is highly recommended not to specify the number of threads, as octopus can probably optimise thread usage better than the user, and it also enables use of specialised multithreaded algorithms which are not available when the thread count is restricted.

When running a multithreaded job, it is also important to consider the sizes of the reference cache and read buffer which are set with the `-x` and `-B` commands respectively. Octopus will always try to respect these buffer sizes; even when using multiple threads. In doing so, a small read buffer size can see each thread being assigned little work, and a small reference cache can lead to cache thrashing. While a larger reference cache size will always improve runtime performance, increasing the read buffer sizes too much can lead to large workloads for each thread which is undesirable as overall thread throughput can decrease. The optimal balance is highly dependent on the data, most critically the depth of coverage. It is recommended the user experiment with different buffer sizes to find an optimal throughput. As a rough guide it is recommended to at least double the default reference cache and read buffer sizes when using multithreading, and quadruple the default sizes when running on a machine with more than 100 cores.

Variant generation

Octopus uses two default candidate variant generators; the raw cigar generator and the local reassembly generator. While it is almost never a good reason to disable the raw cigar generator, it is worth considering the benefit of leaving the assembler enabled, which can bring significant runtime overheads. The main reason to have the assembler enabled is to resolve larger indels and complex variation. As a rule of thumb, the assembler will have less impact with greater read length and decreasing sample diversity. For samples with high diversity, it is recommended to leave the assembler on.

Haplotype generation and phasing

Haplotype generation and phasing are closely related; longer haplotypes will produce longer range variant phasing, and in general, longer haplotypes usually result in more accurate variant calls. Therefore it is important to recognise that adjusting the level of phasing will also impact the quality of variant calling.

By default phasing is medium range. This means octopus will perform haplotype extension when possible, but will stop if too much time is spent in a particular region, or if there are too many haplotypes supported by the data. Experimentation has shown this to provide the best overall calling accuracy while also giving accurate phasing in the vast majority of cases. For human data, the exception is the HLA which may benefit from higher phase levels. In general, increasing the phasing level beyond the default level will not usually improve variant call accuracy, and may decrease it, unless the number of allowed haplotypes is also increased, as the risk of pruning a true haplotype increases with each haplotype extension.

Calling model selection and parametrisation

The most important consideration to make before variant calling is the calling model to use. Using the population calling model on tumour samples will not result in high quality calls (or classified somatic mutations). In most instances, if all relevant sample information is correctly entered to the command line then octopus will automatically select the appropriate calling model. However, there may be exceptions (e.g. tumour only calling), so it is important to check which calling model was invoked when execution begins, or explicitly set the calling model.

Once the appropriate calling model has been chosen, it is vital to consider the parameters specific to that model. For example, the *de novo* mutation rate is specific to the trio calling model. The more accurately the calling model is parametrised, the more accurate the calls will be. It is therefore worth spending time reviewing the parameters.

Some parameters are common to all calling models. Of these, the copy number directives are the most important. In addition to choosing a default organism ploidy, octopus also allows copy number specification of individual chromosomes, so it is important to set correct copy number for sex chromosomes if applicable.

Variant filtering

Variant filtering is used to remove false positive calls that may be introduced due to systematic errors in sequencing or mapping. Ideally, these sources of errors would be fully modelled by the data likelihood model, but capturing all types of error at this stage is extremely difficult. A number of approaches to variant filtering have been proposed, including simple threshold based approaches and sophisticated methods using machine learning. However, all approaches first require defining a set of statistics, or *measures*, that will be used to classify calls as passing or failing in one way or another. The quality of these statistics will ultimately decide the accuracy of variant filtering, regardless of the actual methodology implemented. The default read filter has been chosen to minimise the chance of filtering true positives, whilst eliminating high quality false positives. To achieve very high specificity, it may be necessary to increase the stringency of the filter conditions.

Not all measures available in Octopus are computed during the calling phase, hence some filter expressions require re-access to the read data. The main reason for this is that it may be beneficial to relax the read filtering constraints used for calling compared to filtering. For example, during calling it is usually advisable to filter reads with mapping quality less than 20 as these reads are a common source of false positives. To use them during the calling step would likely increase the false positive rate considerably and increase computation time as more candidates would need to be considered. However, these reads are useful for filtering as they indicate the region where the reads are mapped is likely to contain mapping artefacts.

Octopus currently provides simple threshold based filtering. A number of measures that can be used to define a Boolean *filter expression*. In this version of Octopus, only Boolean *OR* operations are permitted in the filter expression. Note the measure name must appear on the left hand side of each condition in the expression.

Measure name	Default condition	Requires reads	Explanation
AF	< 0.05	Yes	The minimal minor allele frequency in all samples. For germline calling, a low value indicates the variant likely to be sequencing error.
DP	N/A	No	The read depth used for calling is an important consideration as low depth implies low quality genotypes, whilst high depth can result in high quality false positives.
MQD	> 0.6	Yes	The mapping quality divergence is a measure of the difference between the mapping quality distribution between the reads supporting each allele in the called genotype. A large disparity is often indicative of false positives due to mapping error. Octopus computed the Kullback–Leibler divergence of the empirical mapping quality distributions between reads supporting each called allele. A value close to zero indicates low disparity whilst a value closer to one indicates high disparity.
MQ	< 30	Yes	A low mean mapping quality in the region indicates calls may be false positives due to mapping error.
MP	< 20	No	The model posterior is the probability the model Octopus used for calling is true. It is the second best measure to use for filtering after QUAL as it is based on a proper Bayesian model comparison. The model used for comparison varies between calling models, but usually contains an augmented copy number model, which can indicate mapping problems.
QUAL	< 10	No	The call quality is the single best measure to use for filtering as it is derived from a full probability model used during calling. A low value indicates a high probability of false positives.
QD	N/A	No	The quality by depth is the call quality divided by the depth used for calling. This is sometimes useful as some calls may be assigned high quality simply due to high coverage.
SB	> 0.95	Yes	Strand Bias is calculated by comparing the directions of reads supporting each allele in the called genotype. Octopus infers a probability distribution for read direction for each allele and numerically integrates the differences of each pairwise direction bias being over 0.25 for each allele. In other words, the strand bias is the probability the read direction for reads supporting each allele in the called genotype differ by 0.25 or more. A high value therefore indicates the call may be a false positive due to sequencing errors.

It is important to consider the increased computation burden that some measures introduced. In particular any measures that require access to read data will increase memory consumption and runtime. Some measures

also require computation of alignments to called genotypes to calculate supporting reads (e.g. for strand bias), which again increases computation time.

EXAMPLES

This section contains some common use-case examples. The section is not intended as a replacement for the command line reference and performance optimisation sections, which should be consulted to tune parameters for specific user requirements. Note all of the examples in this section use the default output mode (standard output) for brevity, to write to a file just add the `--output; -o` command.

Calling germline variants in a single sample

As previously described, octopus has two distinct models for germline variant calling - one for a single individual and another for populations. Fortunately there is no concern for the user as the appropriate model is selected automatically:

```
$ octopus -R human.fa -I NA12878.bam
```

Assuming the file `NA12878.bam` contains a single sample, this will use the *individual* calling model. Octopus does not care how many samples are actually in a read file, so if the input read file contains multiple samples but only a single sample is required for analysis, the name of the sample is required as input:

```
$ octopus -R human.fa -I multi_sample.bam --samples NA12878
```

Calling variants in a targeted exome panel

All octopus calling models can be supplied with a list of target intervals to analyse, for a small number of regions the option `--regions; -T` can be used:

```
$ octopus -R human.fa -I NA12878.bam -T 22:35,799,116-35,799,685
```

This option can be used multiple times, or can be supplied with a space separated list of arguments. However, for longer target interval lists it may be easier to create a file which lists the regions (one per line) and pass this to octopus using the `--regions-file; -t` option:

```
$ octopus -R human.fa -I NA12878.bam -t exome-panel.txt
```

Note these options can be used in conjunction, and there is no need to worry about duplicates or overlaps - octopus will resolve this internally.

Ignoring decoy contigs from a whole genome run

There is a useful option, `--skip-regions; -K`, that serves as the converse of the `--regions` command; it informs octopus **not** to analyse the given options. The main utility of this is for ignoring decoy contigs or centromeres in whole genome runs. There is a homologous command, `--skip-regions-file; -k`, which takes a path to a file containing regions to ignore:

```
$ octopus -R human.fa -I NA12878.bam -k human-decoy.txt
```

It is possible to use all the region specific commands in conjunction to get fine grain control over which regions to call.

Calling germline variants in a population (in development)

The population model is the default for more than a single sample, so just supply a list of samples:

```
$ octopus -R human.fa -I NA12878.bam NA12891.bam
```

For larger sample sets, it is usually better to have the read paths in a file:

```
$ octopus -R human.fa -i reads.txt
```

Calling *de novo* mutations in a trio

To call germline and *de novo* mutation in a trio, just specify `--maternal-sample; -M` and `--paternal-sample; -F`:

```
$ octopus -R human.fa -i ceu_trio.txt -M NA12892 -F NA12892
```

The child is automatically deduced. The trio can also be specified with a PED file:

```
$ octopus -R human.fa -i ceu_trio.txt --pedigree ceu_trio.ped
```

Calling somatic mutations in a tumour-normal pair

To call germline and somatic variants in tumour samples, supply either a normal sample:

```
$ octopus -R human.fa -I normal.bam tumour.bam -N normal
```

If a normal sample is unavailable, tumour only calling can be invoked by explicitly selecting the cancer calling model:

```
$ octopus -R human.fa -I tumour1.bam tumour2.bam -C cancer
```

HLA genotyping

Octopus is able to call very long haplotypes, especially in variant dense regions, which makes it an ideal tool for calling HLA haplotypes. By default octopus will not make maximally long haplotypes - and therefore phase regions - due to the computational complexity involved in such optimisation, and the diminishing return of very long haplotypes. But in the HLA, longer haplotypes are desired, which can be achieved using the `--phasing-level; -l` command:

```
$ octopus -R human.fa -I NA12878.bam -t hla-regions.txt -l aggressive
```

It may also be beneficial to increase the default value of `--max-haplotypes` to 256 or 512.

Calling variants in haploid organism

The default parameters are set with human sequence data in mind, for non-human samples it is recommended to adjust the options for the organism being analysed. For haploid organisms such as bacteria and viruses, the most important parameter to change is `--organism-ploidy; -P` which sets the default

ploidy to use. Depending on the organism, it may also be important to adjust the variant priors: `--snp-heterozygosity` and `--indel-heterozygosity`:

```
$ octopus -R ecoli.fa -I ecoli.bam -P 1 --snp_heterozygosity 0.01
```

Running in multithread mode

By default all octopus runs execute using a single thread, but it is trivial to use multiple threads using the `--threads` command:

```
$ octopus -R human.fa -I NA12878.bam --threads
```

This is the recommended approach to multithreading with octopus, but the command also takes an optional number of threads to use, which must be specified immediately after the command:

```
$ octopus -R human.fa -I NA12878.bam --threads=4
```

The former form is recommended because it allows octopus to optimise thread usage, and also enables the use of specific multithreaded algorithms.

Using a configuration file

Octopus allows all command line options to be specified using a configuration file, which some users may prefer is the same configuration us used often. The configuration file is just a text file with each line containing a `option=value` pair:

```
$ octopus --config my_octopus_config.txt
```

TROUBLESHOOTING

BUILDING

Why are the requirements so strict?

Octopus uses some advanced features of the latest official C++ standard (C++14) and therefore requires a mature C++ compiler. The other requirements have been set to recent versions because these are the versions used to develop and test octopus, and we cannot guarantee that earlier versions will perform as well. This also reduces the burden on future releases to be backwards compatible with ageing tools, which allows focus on new features and improvements.

CMake chooses a bad compiler

As the previous issue explains, some compilers advertised as being C++14 compliant are not due to compiler bugs. Unfortunately, CMake cannot recognise this and will happily select a buggy compiler if you have one in a standard location on your system. If you have installed another working compiler in a non-standard location, you need to tell CMake how to find it with the `CMAKE_CXX_COMPILER` command:

```
$ cmake -D CMAKE_CXX_COMPILER=/path/to/compiler/cpp ..
```

Or if you're using the Python install script (recommended):

```
$ ./install.py --compiler=/path/to/compiler/cpp
```

Compilation fails

Octopus uses advanced C++14, and therefore requires a robust C++14 compiler and standard library implementation. Many of the compiler versions advertised as being C++14 compliant have bugs that prevent compilation. The only compilers that have currently been shown to work are Clang 3.8 and GCC 6.2.

Linking fails

Ensure the correct compiler driver was selected: using Clang this means selecting `clang++` and not `clang`:

```
$ ./install.py --compiler=/path/to/clang/bin/clang++
```

Similarly for GCC use `g++` rather than `gcc`:

```
$ ./install.py --compiler=/path/to/gcc/bin/g++
```

If the issue is specifically with linking against Boost then see the next issue.

Boost libraries fail to link

First ensure that Boost was built using the same compiler used to compile octopus - using different compilers will often cause linking problems.

Second, if the required Boost libraries are not installed in a standard location on your system, and you have not set appropriate environment variables, you may see an error message when running octopus (even after a successful build). To resolve this you just need to set the appropriate environment variable:

```
$ export LD_LIBRARY_PATH=/path/to/boost/lib
```

before executing octopus.

Compilation has lots of #pragma warnings

This is a Boost issue that was introduced in Boost 1.60, and should be fixed in future releases of Boost. This issue does not affect octopus.

RUNTIME

Segmentation fault

Sorry! Please see the advise under 'reporting bugs'.

Execution is slow

The best way to improve runtime is to use multithreading with the `--threads` command. If this does not offer adequate performance then consider switching off the assembler with the `-a` command or try the command `--fast`. Please see the section Execution time under performance optimisation.

Execution delays after initialising calling components in threaded mode

This is due to the way tasks are distributed to threads: the challenge is to split the genome into chunks (or tasks) that are *independent* in the sense that the union of the calls produced by calling each independently is the same as calling them together. Clearly this is always true for different contigs, but it is non-trivial to detect independent tasks for a single contig. In addition, if tasks are too small or too large, runtime performance is not optimal. Octopus solves this problem by focusing on creating tasks that will lead good thread usage, while relying on a conflict detection algorithm to find non-independent tasks *after* calling has completed. This algorithm requires all tasks for a given contig to be known before any resolutions can be made, and therefore all tasks for a contig must be generated before calling can begin.

Run hangs in decoy contigs

Decoy contigs often have very high coverage (above 20,000X) and therefore the downsampling will likely occur. The downsampling algorithm octopus implements is non-trivial as it attempts to keep an even coverage across the downsampled region to avoid any bias. Unfortunately in very high coverage regions this can be very slow. We are working on improving this, but for now we suggest either increasing the downsampling thresholds, or removing these decoy contigs from analysis completely (using the option `--skip-regions (-file)`).

BEHAVIOUR

No calls are reported

Try turning off read filters with the command `--disable-read-filtering`. If this results in calls then one or more of the filters is probably over-filtering, which usually implies a quirk of the read mapper that octopus does not recognise. In such cases it is best to find which filter is triggering (the `--debug` command can be used to help) and disable it.

Regions are skipped because of too many haplotypes

This warning can occur in very complex regions and is not normally a problem as it is usually reflective of bad read mappings. However, it is possible to force octopus to call these regions by increasing the `--max-haplotypes` and `--haplotype-overflow` command line options.

A call changes when a different input region is given

Haplotype based variant callers gain power by jointly calling adjacent alleles, hence the result for one position is dependent on other nearby positions. If two target regions differ then different haplotypes may be proposed for each, which can affect the overall result for a position even if it is present in both target regions, and even if no other calls result from the difference. Octopus tries to avoid this problem as much as possible by considering regions beyond those requested, but this cannot eliminate the problem entirely - the only complete solution is to call entire contigs.

Why doesn't octopus report genotype likelihoods?

Octopus calculates the likelihood of the reads given a *haplotype* while records in the VCF file are *alleles*. It is non-trivial to calculate the genotype likelihood w.r.t an allele given the genotype likelihoods w.r.t haplotypes as the calculation must also condition on all other alleles used to compute the haplotypes likelihoods. This is not the case for posteriors, which can be marginalised in a trivial way, as all required information regarding *other* alleles is already present in the single posterior value.

While we appreciate users like to see genotype likelihoods. We feel we can offer users more accurate results by providing flexible priors, and reporting exact genotype posteriors in the final output.

Why do octopus VCF files contain * and .?

Octopus uses advanced parts of the VCF specification to archive a consistent genotyping over all samples. The problem appears because primarily because multi-allelic sites must either be represented on a single line or split into multiple records. For shorter alleles the former option is fine, but the latter option is better for longer alleles (otherwise records containing 100+ bases to represent a SNP occur). The principle problem with splitting records like this is that inconsistencies can arise if the records are treated independently, for example, consider a simple example:

#CHROM	POS	ID	REF	ALT	QUAL	FILTER	INFO	FORMAT	NA12878
1	100	.	A	C,G	.	.	.	GT	1 2

Now suppose we split the record into two separate records:

#CHROM	POS	ID	REF	ALT	QUAL	FILTER	INFO	FORMAT	NA12878
1	100	.	A	C	.	.	.	GT	1 0
1	100	.	A	G	.	.	.	GT	0 1

There is an inconsistency here! The first record states the genotype is C|A while the second claims it is A|G! The problem is that the reference is overrepresented in the genotype call. What we should be saying here is that the genotype cannot be fully represented due to a conflict with another call, which is exactly what octopus does!

#CHROM	POS	ID	REF	ALT	QUAL	FILTER	INFO	FORMAT	NA12878
1	100	.	A	C	.	.	.	GT	1 . .
1	100	.	A	G	.	.	.	GT	. 1

Note: octopus would actually represent this call as a multi-allelic record (as in the first case), this example is just intended to demonstrate the behaviour.

Now for the “*” which can appear as an ALT allele. This is somewhat like the dot used to indicate an interaction, but specifically denotes an overlap with a deletion.

Unfortunately, many other variant callers - and therefore downstream analysis tools - do not output consistent calls. In order to force octopus to output an additional potentially inconsistent VCF file which may work better with downstream tools, just add the option `--legacy`.

SNP accuracy improves in fast mode

In fast mode octopus turns off haplotype extension, which means there is no haplotype posterior based filtering. Haplotype extension is usually good; it can help resolve complex regions and allows long range phasing. However, in regions where the data cannot be modelled correctly (e.g. if many reads are miss-mapped), it can lead to significant portions of the overall posterior distribution being removed. This causes false positive calls to be given far higher posterior (and hence QUAL) than if no extension was used. This is especially true for SNPs.

Many of these high quality false positives will be filtered with the default model filters, but it is possible for some to pass these filters if too many reasonable haplotypes are removed before the model selection is applied. We are working on ways to improve this, one idea is to do a double pass of any complex regions to improve resolution. If this is causing a significant problem for you now, you can help alleviate the issue by increasing the posterior filter threshold (`--haplotype-extension-threshold`), or simply turn off haplotype extension by setting `--phasing-level` to `minimal`.

Calling performance is worse with assembler

The assembler often proposes many complex candidate variants which increases the number of haplotypes considerably. This can force octopus to rely on its haplotype filtering algorithms more than without the assembler, and while these algorithms often perform well, they are not foolproof. In addition, many gold standard reference sets will not contain real complex variants of the type the assembler is able to propose, as these are challenging to validate, and many variant callers are not even capable of calling such variants. It is advised the user interpret such results with a degree of caution.

CONTACT

- Author and maintainer: Daniel Cooke (dcooke@well.ox.ac.uk).
- Author: Gerton Lunter (gerton.lunter@well.ox.ac.uk).

APPENDIX

INSTALLING REQUIREMENTS

OS X

On OS X, Clang is recommended, and all requirements can be installed with the package manager Homebrew:

```
$ brew update
$ brew install git
$ brew install --with-clang llvm
$ brew install boost
$ brew install cmake
$ brew install python3
$ brew tap homebrew/science
$ brew install htplib
```

If you already have any of these packages installed on your system with Homebrew the command will fail, but you can update to the latest version by using `brew upgrade` instead of `brew install`.

Ubuntu

On Ubuntu, most requirements can be obtained with `apt-get`. Clang 3.8 is recommended as GCC 6.2 has a bug which reduces octopuses runtime performance. To obtain the requirements execute:

```
$ sudo apt-get update && sudo apt-get upgrade
$ sudo apt-get install clang-3.8
$ sudo apt-get install libstdc++6
$ sudo apt-get install libboost-all-dev
$ sudo apt-get install cmake
$ sudo apt-get install git-all
$ sudo apt-get install python3
```

As only htplib 1.2.1 is available using `apt-get`, htplib 1.3.1 must be installed manually:

```
$ sudo apt-get install autoconf
$ git clone https://github.com/samtools/htslib.git
$ cd htslib
$ autoheader
$ autoconf
$ ./configure
$ make && sudo make install
```

VARIANT GENERATION

Genotype calls are made by generating a set of candidate variants and weighing up the evidence for each. The space of all possible variants is infinite, so heuristics must be employed to generate the candidate set. The performance of the variant generator is critical as the final calls will be a subset of the generated candidate variant set.

Octopus can make use of multiple independent variant generators and take the union of the result of each. The two default variant generators are the **raw cigar generator** and the **local reassembly generator**. The raw cigar generator proposes a candidate whenever a mismatch, as indicated in a reads cigar string, satisfies some condition. The default condition used is dependent on the calling model selected, but can be modified by the user. The local reassembly generator constructs a De Bruin graph of all reads in a genomic interval, threaded around the reference sequence, and then extracts paths through the graph which deviate from the single reference path.

Currently the only other non-default generator is the **source file generator** which simply extracts previously made calls from a VCF format file. We are working on a generator that will extract known variants from online databases.

HAPLOTYPE GENERATION

The haplotype generator takes the set of candidate variants generated by the variant generators and constructs sets of candidate haplotypes. The advantage of having these two stages independent is that haplotypes can then be dynamically manipulated after construction. In particular, it allows a haplotype to be extended conditionally on the posterior probability that the haplotype is present in the union of all sample genotypes.

PHASING

Phasing takes place after variants have been called - the algorithm uses the regions spanned by the calls and the inferred genotype posterior distribution to find optimal phase regions. It proceeds by arranging the called regions into groups such that the entropy of marginal posterior distributions between groups is maximised. In other words, if the posterior distribution assigns similar mass to different phasing of the same variant calls,

then the true phasing is unknown, but if one particular phasing carries most of the posterior mass, then there is strong evidence that phasing is the true physical phasing.

VARIANT CALL FILTERING

The octopus calling stage is designed to be as sensitive as possible. However, increased sensitivity usually comes at a cost of decreased specificity. Most false positive calls fall into three categories (discounting software errors -bugs - and user errors):

- Mapping based errors: the supporting a false positive call are mapped to the wrong location.
- Sequencing errors: there is sufficient sequencing error (usually indicative of systematic error) to make a false positive call.
- Model errors: the models used to call variants does not fit the target data.

Of these, mapping based errors are the most difficult to resolve as there is no consistent model that would explain miss-mapped reads better than correctly-mapped reads.

GLOSSARY

Allele A particular instance of a genomic region, that is, a nucleotide sequence and a genomic region.

Genomic region/interval A coordinate range with respect to a reference genome. All genomic regions must specify the contig (chromosome) and begin and end positions in the reference.

Genotype A collection of haplotypes or alleles of known cardinality.

Haplotype An ordered set of alleles which occur on the *same* chromosome in the sample.